

mhshlgisj

April 27, 2024

## 1 Objective

- Use a business dataset to formulate conclusions
- Create visual representations of data collected through data wrangling and data filtration
- Interpret datasets using data retrieval techniques to collect them

## 2 Problem Area

When users look for a place to stay, they may look for Airbnb's or hotels. In this case, I will be looking at the dataset regarding New York City's Airbnb Options. Through data retrieval techniques, I will provide information regarding a house's rental price and its area. I will also look at the Airbnb properties of the Airbnb hosts.

By collecting this information, both customers and hosts can benefit from this. Customers can use this data to filter a location that is most cost-effective for them, providing them with informed decisions on where to stay. They will also be able to see which host offers locations all year round and which do not. Hosts, on the other hand, can see which locations yield the highest return for their money.

These findings are part of Airbnb's public dataset.

## 3 Understanding the Dataset

The dataset I will be looking at can be found through Kaggle's collection of datasets. (<https://www.kaggle.com/datasets/dgomonov/new-york-city-airbnb-open-data>)

This dataset involves an official release of New York City's Airbnb open data from 2019 taken directly from the Airbnb website.

I will be looking primarily at the **price, host name, neighborhood, and availability**.

Understanding rental price trends will provide valuable information to the hosts and guests by filtering information such as cost and area to those interested in the Airbnb market.

## 4 Queries

By answering and plotting these queries, I will provide an analysis of how each query can be interpreted.

1. Top 3 neighborhoods that have the highest average nightly rental price

2. Top 3 neighborhoods that have the lowest average nightly rental price
3. Top 10 hosts who have the lowest nightly rental price
4. List the name's of 10 hosts who have had at least one of their Airbnb listings available to rent all year
5. List the name's of 10 hosts who have had at least one of their Airbnb listings not available to rent all year
6. Top 10 hosts who have the highest nightly rental price

I will begin by first importing the required libraries used for retrieving and plotting data.

```
[1]: import pandas as pd
import matplotlib.pyplot as plt
```

```
[2]: data = pd.read_csv('AB_NYC_2019.csv')
```

## 5 Pre-Processing/Data Cleaning

Since I only need to work with the columns pertaining to host name, neighborhood, price, and availability. I will only focus on those by gathering the element number correlated to their respective columns. This will provide me of a subset of the relevant data.

```
[3]: data.columns
```

```
[3]: Index(['id', 'name', 'host_id', 'host_name', 'neighbourhood_group',
          'neighbourhood', 'latitude', 'longitude', 'room_type', 'price',
          'minimum_nights', 'number_of_reviews', 'last_review',
          'reviews_per_month', 'calculated_host_listings_count',
          'availability_365'],
          dtype='object')
```

```
[4]: relevantData = data.iloc[:, [3, 5, 9, 15]].copy()
relevantData
```

```
[4]:
```

	host_name	neighbourhood	price	availability_365
0	John	Kensington	149	365
1	Jennifer	Midtown	225	355
2	Elisabeth	Harlem	150	365
3	LisaRoxanne	Clinton Hill	89	194
4	Laura	East Harlem	80	0
...	...	...	...	...
48890	Sabrina	Bedford-Stuyvesant	70	9
48891	Marisol	Bushwick	40	36
48892	Ilgar & Aysel	Harlem	115	27
48893	Taz	Hell's Kitchen	55	2
48894	Christophe	Hell's Kitchen	90	23

```
[48895 rows x 4 columns]
```

Here I check for any null values in my relevant data using `.isnull()` and drop any null values I find after locating which column is associated with the null values.

```
[5]: relevantData.isnull().sum()
```

```
[5]: host_name          21  
     neighbourhood      0  
     price              0  
     availability_365    0  
     dtype: int64
```

```
[6]: relevantData.dropna(how = 'all', subset = ['host_name'], inplace = True)  
     relevantData.isnull().sum()
```

```
[6]: host_name          0  
     neighbourhood      0  
     price              0  
     availability_365    0  
     dtype: int64
```

After dropping null values, I also delete any duplicate rows that can provide inconsistencies in the data processing process.

```
[7]: relevantData.duplicated().sum()
```

```
[7]: 895
```

```
[8]: relevantData.drop_duplicates(inplace = True)  
     relevantData.duplicated().sum()
```

```
[8]: 0
```

To see what type of information the Airbnb dataset provides me with, I will print then data types for each column in the relevant data table I have created.

```
[9]: relevantData.dtypes
```

```
[9]: host_name          object  
     neighbourhood      object  
     price              int64  
     availability_365    int64  
     dtype: object
```

Since the average price corresponding to the neighborhood will be used in future queries, I will add a column of average rental prices in addition to the data I have previously sliced.

```
[10]: neighborhood_avg_price = relevantData.groupby('neighbourhood')['price'].  
      ↪transform('mean')  
      relevantData['avg_price'] = neighborhood_avg_price
```

```
relevantData
```

```
[10]:
```

	host_name	neighbourhood	price	availability_365	avg_price
0	John	Kensington	149	365	93.195402
1	Jennifer	Midtown	225	355	274.896480
2	Elisabeth	Harlem	150	365	119.500764
3	LisaRoxanne	Clinton Hill	89	194	181.391228
4	Laura	East Harlem	80	0	133.812500
...	...	...	...	...	...
48890	Sabrina	Bedford-Stuyvesant	70	9	108.525130
48891	Marisol	Bushwick	40	36	85.446013
48892	Ilgar & Aysel	Harlem	115	27	119.500764
48893	Taz	Hell's Kitchen	55	2	203.776596
48894	Christophe	Hell's Kitchen	90	23	203.776596

```
[47979 rows x 5 columns]
```

## 6 1. Top 3 neighborhoods that have the highest average nightly rental price

Using `sort_values()`, I will obtain the average price and order the dataset by descending order to gather the average prices starting from the highest to lowest.

By using `drop_duplicates()`, I will drop any duplicates of the same neighborhood to gather the top 3 neighborhoods instead of the top 3 average prices only since there can be multiple Airbnb's in one neighborhood.

The use of `.head(3)` will gather the top 3 results of the table.

```
[11]: top_3_neighborhoods_avg = relevantData.sort_values(by = 'avg_price', ascending_
      ↪ = False).drop_duplicates(subset = 'neighbourhood').head(3)
      top_3_neighborhoods_avg
```

```
[11]:
```

	host_name	neighbourhood	price	availability_365	avg_price
25386	Mark	Fort Wadsworth	800	365	800.000000
3049	Donna	Woodrow	700	0	700.000000
46252	Ej	Tribeca	600	80	493.028409

By creating a bar graph, we can plot the top 3 neighborhoods associated with the highest average nightly rental price and see the price associated with it.

```
[12]: avgPrice = top_3_neighborhoods_avg['avg_price']
      neighborhoods = list(top_3_neighborhoods_avg['neighbourhood'])


      fig,ax = plt.subplots(nrows = 1, ncols = 1)

      ax.bar(neighborhoods, avgPrice)
```

```
plt.xlabel('Neighborhoods', fontsize = 14)
plt.ylabel('Average Price', fontsize = 14)
plt.xticks(rotation = 'vertical', fontsize = 17)
plt.title('Top 3 Neighborhoods that have the Highest Average Nightly Rental_
↪Price', fontsize = 25)

ax.grid()
plt.show()
```

Top 3 Neighborhoods that have the Highest Average Nightly Rental Price



For Airbnb customers, these neighborhoods may be areas that customers may want to deter from. For hosts, these neighborhoods may be areas to purchase properties to sell for rental.

## 7 2. Top 3 neighborhoods that have the lowest average nightly rental price

In the code snippet below, the average price is sorted this time without stating ascending to be false. This will make ascending true by default, and the `drop_duplicates()` is implemented to remove the same neighborhood in case another row has a duplicate. Since we want to see the top 3 neighborhoods and not lowest average nightly rental price only, it is necessary.

`.head()` is used to obtain the top 3.

```
[13]: top_3_neighborhoods_lowest_avg = relevantData.sort_values(by = 'avg_price').
↪drop_duplicates(subset = 'neighbourhood').head(3)
top_3_neighborhoods_lowest_avg
```

```
[13]:
```

	host_name	neighbourhood	price	availability_365	avg_price
37044	Xiomara	Bull's Head	80	362	47.333333
29591	Schomberg	Hunts Point	45	0	50.500000
22708	Emma	Tremont	38	327	51.545455

This code block plots the top 3 neighborhoods with the least average nightly rental price.

```
[14]: avgPrice = top_3_neighborhoods_lowest_avg['avg_price']
neighborhoods = list(top_3_neighborhoods_lowest_avg['neighbourhood'])

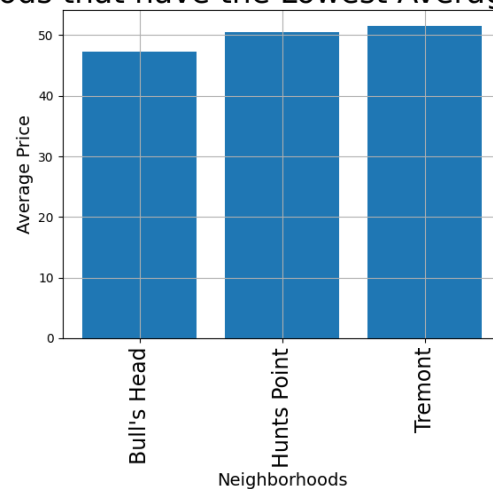
fig, ax = plt.subplots(nrows = 1, ncols = 1)

ax.bar(neighborhoods, avgPrice)

plt.xlabel('Neighborhoods', fontsize = 14)
plt.ylabel('Average Price', fontsize = 14)
plt.xticks(rotation = 'vertical', fontsize = 17)
plt.title('Top 3 Neighborhoods that have the Lowest Average Nightly Rental_
↪Price', fontsize = 25)

ax.grid()
plt.show()
```

Top 3 Neighborhoods that have the Lowest Average Nightly Rental Price



For Airbnb customers, these neighborhoods may be areas that customers may be attracted to since they are more cost-effective. For hosts, these neighborhoods may be areas to deter away from since they yield lower rental prices.

### 8 3. Top 10 hosts who have the lowest nightly rental price

I collect the top 10 expensive nightly rental prices here by sorting the table by its price in ascending order. I then drop any duplicates of the host name in case a host has multiple properties of the cheapest prices.

I then use `.head(10)` to output the top 10 hosts who have the lowest nightly rental price.

```
[15]: top_10_host = relevantData.sort_values(by='price').drop_duplicates(subset =  
      ↪ 'host_name').head(10)  
top_10_host
```

```
[15]:
```

	host_name	neighbourhood	price	availability_365	\
26259	Qiuchi	Murray Hill	0	0	
25753	Lauren	Greenpoint	0	0	
26866	Sergii	Bushwick	0	139	
25778	Aymeric	Williamsburg	0	73	
25794	Adeyemi	Bedford-Stuyvesant	0	176	
23161	Kimberly	Bedford-Stuyvesant	0	28	
25433	Anisha	East Morrisania	0	127	
25634	Martial Loft	Bushwick	0	0	
47218	Julio	Bushwick	10	0	
27972	Vishanti & Jeremy	Greenpoint	10	32	

	avg_price
26259	224.182018
25753	143.850788
26866	85.446013
25778	144.659802
25794	108.525130
23161	108.525130
25433	85.000000
25634	85.446013
47218	85.446013
27972	143.850788


In the code snippet below, I output the result in a bar chart where the top 10 hosts who have the lowest nightly rental price are shown with their corresponding price.

```
[16]: price = top_10_host['price']  
host = list(top_10_host['host_name'])  
  
fig,ax = plt.subplots(nrows = 1, ncols = 1)  
  
ax.bar(host, price)  
  
plt.xlabel('Host', fontsize = 14)  
plt.ylabel('Price', fontsize = 14)
```

```
plt.xticks(rotation = 'vertical', fontsize = 17)
plt.title('Top 10 Hosts who have the Lowest Nightly Rental Price', fontsize = 25)

ax.grid()
plt.show()
```

## Top 10 Hosts who have the Lowest Nightly Rental Price



As shown in the bar chart, the prices and hosts are given for the lowest prices. For customers, this can be used as a filter to look at any other listings the host may have since they have cost-effective prices. For hosts, this can help bring business to them since they will stand out as having cheaper properties.

## 9 4. List the name's of 10 hosts who have had at least one of their Airbnb listings available to rent every day of the year.

In the code snippet, I use the relevantData table I created to look for the number '365' in the availability\_365 column. By setting this condition in place, the table will output only rows with that condition being met, which indicates that the Airbnb listing has been available to rent every day. I drop duplicates of the host name since every row must have a unique host name.

Since order is not relevant here, I will not be sorting by ascending or descending.

I then use .head(10) to return 10 rows the table outputted.



```
[17]: all_year_availability = relevantData[relevantData['availability_365'] == 365].
      ↪drop_duplicates(subset = 'host_name').head(10)
all_year_availability
```

```
[17]:
```

	host_name	neighbourhood	price	availability_365	avg_price
0	John	Kensington	149	365	93.195402
2	Elisabeth	Harlem	150	365	119.500764
36	Vt	Bedford-Stuyvesant	35	365	108.525130
38	Harriet	Flatbush	150	365	92.321370
97	Elliott	Harlem	89	365	119.500764
139	Sahr	Clinton Hill	135	365	181.391228
164	Sylvia	Upper East Side	250	365	189.402715
181	Mark	Ridgewood	350	365	77.654589
196	Jullett	Jamaica	55	365	94.740909
204	Douglas	Harlem	200	365	119.500764

The code snippet below shwos a bar graph of the hosts who had their listing available to rent 365 days of the year.

```
[18]: availability = all_year_availability['availability_365']
host = list(all_year_availability['host_name'])

fig,ax = plt.subplots(nrows = 1, ncols = 1)


ax.bar(host, availability)

plt.xlabel('Host', fontsize = 10)
plt.ylabel('Availability Days', fontsize = 10)
plt.xticks(rotation = 'vertical', fontsize = 12)
plt.title('10 Hosts who have had at Least One of Their Airbnb Listing Available,
      ↪to Rent All Year', fontsize = 15)

plt.ylim(0,365)

ax.grid()
plt.show()
```

10 Hosts who have had at Least One of Their Airbnb Listing Available to Rent All Year



From the bar graph, you can see at least 10 hosts who have had their listing available 365 days of the year. Customers can use this information to see which listings are not in high demand. They may also use it to filter out for available listings, so this will narrow their search and lead to a better experience. Hosts may benefit from this since it will bring attention to customers who are looking for a low demand listing, and potentially cheaper rental price.

## 10 5. List the name's of 10 hosts who have had at least one of their Airbnb listings not available to rent all year.

In the code snippet, I use the `relevantData` table I created to look for the number '0' in the `availability_365` column. By setting this condition in place, the table will output only rows with that condition being met, which indicates that the Airbnb has not been available to rent at all in one year. I drop duplicates of the host name since every row must have a unique host name.

Since order is not relevant here, I will not be sorting by ascending or descending.

I then use `.head(10)` to return 10 rows the table outputted.

```
[19]: no_availability = relevantData[relevantData['availability_365'] == 0].
      ↪drop_duplicates(subset = 'host_name').head(10)
      no_availability
```

```
[19]:
```

	host_name	neighbourhood	price	availability_365	avg_price
4	Laura	East Harlem	80	0	133.812500
6	Garon	Bedford-Stuyvesant	60	0	108.525130
8	MaryEllen	Upper West Side	79	0	209.478619
14	Alina	West Village	120	0	268.448231

20	Chaya	Williamsburg	299	0	144.659802
26	Claude & Sophie	Inwood	80	0	88.898785
48	Jennifer	Bedford-Stuyvesant	115	0	108.525130
66	Sara	Park Slope	225	0	176.542574
88	Pas	East Village	50	0	186.730811
94	Christiana	Williamsburg	100	0	144.659802

The code snippet below shows a bar graph of the hosts who had their listing available to rent 0 days of the year.

```
[20]: availability = no_availability['availability_365']
host = list(no_availability['host_name'])

fig, ax = plt.subplots(nrows = 1, ncols = 1)


ax.bar(host, availability)

plt.xlabel('Host', fontsize = 10)
plt.ylabel('Availability Days', fontsize = 10)
plt.xticks(rotation = 'vertical', fontsize = 12)
plt.title('10 Hosts who have had at Least One of Their Airbnb Listing not
↳ Available to Rent All Year', fontsize = 15)

plt.ylim(0,365)

ax.grid()
plt.show()
```

10 Hosts who have had at Least One of Their Airbnb Listing not Available to Rent All Year



From the bar graph, you can see at least 10 hosts who have had their listing available 365 days of the year. Customers can use this information to see which listings are in high demand. They may also use it to filter out any Airbnb's that are not available, so this will narrow their search and lead to a better experience. Customers may also take note of how high demand a certain Airbnb is, which may lead to them booking it earlier. Hosts may benefit from this since they will be able to see their competition and see what a successful Airbnb looks like.

## 11 6. Top 10 hosts who have the highest nightly rental price

The code sorts values by price in a descending order by setting ascending = false. It also drops duplicates since every host name should be unique.

.head(10) is used to grab the top 10 hosts.

```
[21]: top_10_host = relevantData.sort_values(by='price', ascending=False).
      ↪ drop_duplicates(subset = 'host_name').head(10)
      top_10_host
```

```
[21]:
```

	host_name	neighbourhood	price	availability_365	avg_price
9151	Kathrine	Astoria	10000	0	117.495536
29238	Jelena	Upper West Side	10000	83	209.478619
17692	Erin	Greenpoint	10000	0	143.850788
40433	Matt	Lower East Side	9999	365	188.828090
6530	Olson	East Harlem	9999	0	133.812500
12342	Amy	Lower East Side	9999	83	188.828090
30268	Rum	Tribeca	8500	251	493.028409
4377	Jessica	Clinton Hill	8000	365	181.391228
29662	Sally	Upper East Side	7703	146	189.402715
45666	Sandra	East Flatbush	7500	179	104.881391

The code snippet below graphs the top 10 hosts who have the highest nightly rental price in a bar chart.

```
[22]: price = top_10_host['price']
      host = list(top_10_host['host_name'])


      fig,ax = plt.subplots(nrows = 1, ncols = 1)

      ax.bar(host, price)

      plt.xlabel('Host', fontsize = 14)
      plt.ylabel('Price', fontsize = 14)
      plt.xticks(rotation = 'vertical', fontsize = 17)
      plt.title('Top 10 Hosts who have the Highest Nightly Rental Price', fontsize = 25)
      ↪
```

```
ax.grid()
plt.show()
```

## Top 10 Hosts who have the Highest Nightly Rental Price



Customers can use this bar chart to see which hosts to potentially avoid or look at if they seek a potentially luxurious place to rent. Hosts can benefit from this because they can see what an expensive Airbnb entails and see their competition in that price area.

## 12 Conclusion

Overall, by analyzing the New York City's Airbnb dataset, we can provide valuable insights on how customers and benefits can benefit from them. Customers can use the information as a way to filter their needs such as sorting by cheaper or expensive rentals. They will be able to filter which neighborhoods yield a higher price on average, and which yield a smaller price. They will also be able to see which Airbnb's are usually more readily available, which may affect their booking date.

Hosts on the Airbnb website hosting rental properties in New York City may also benefit because they also gain valuable insights. Hosts can see which locations yield high-return results when it comes to money. By seeing which locations are more expensive or cheaper on average, they may adjust their pricing accordingly to stay competitive with their fellow hosts.

By plotting these results, I was able to provide the results visually to make it easier to understand the results and identify any trends.